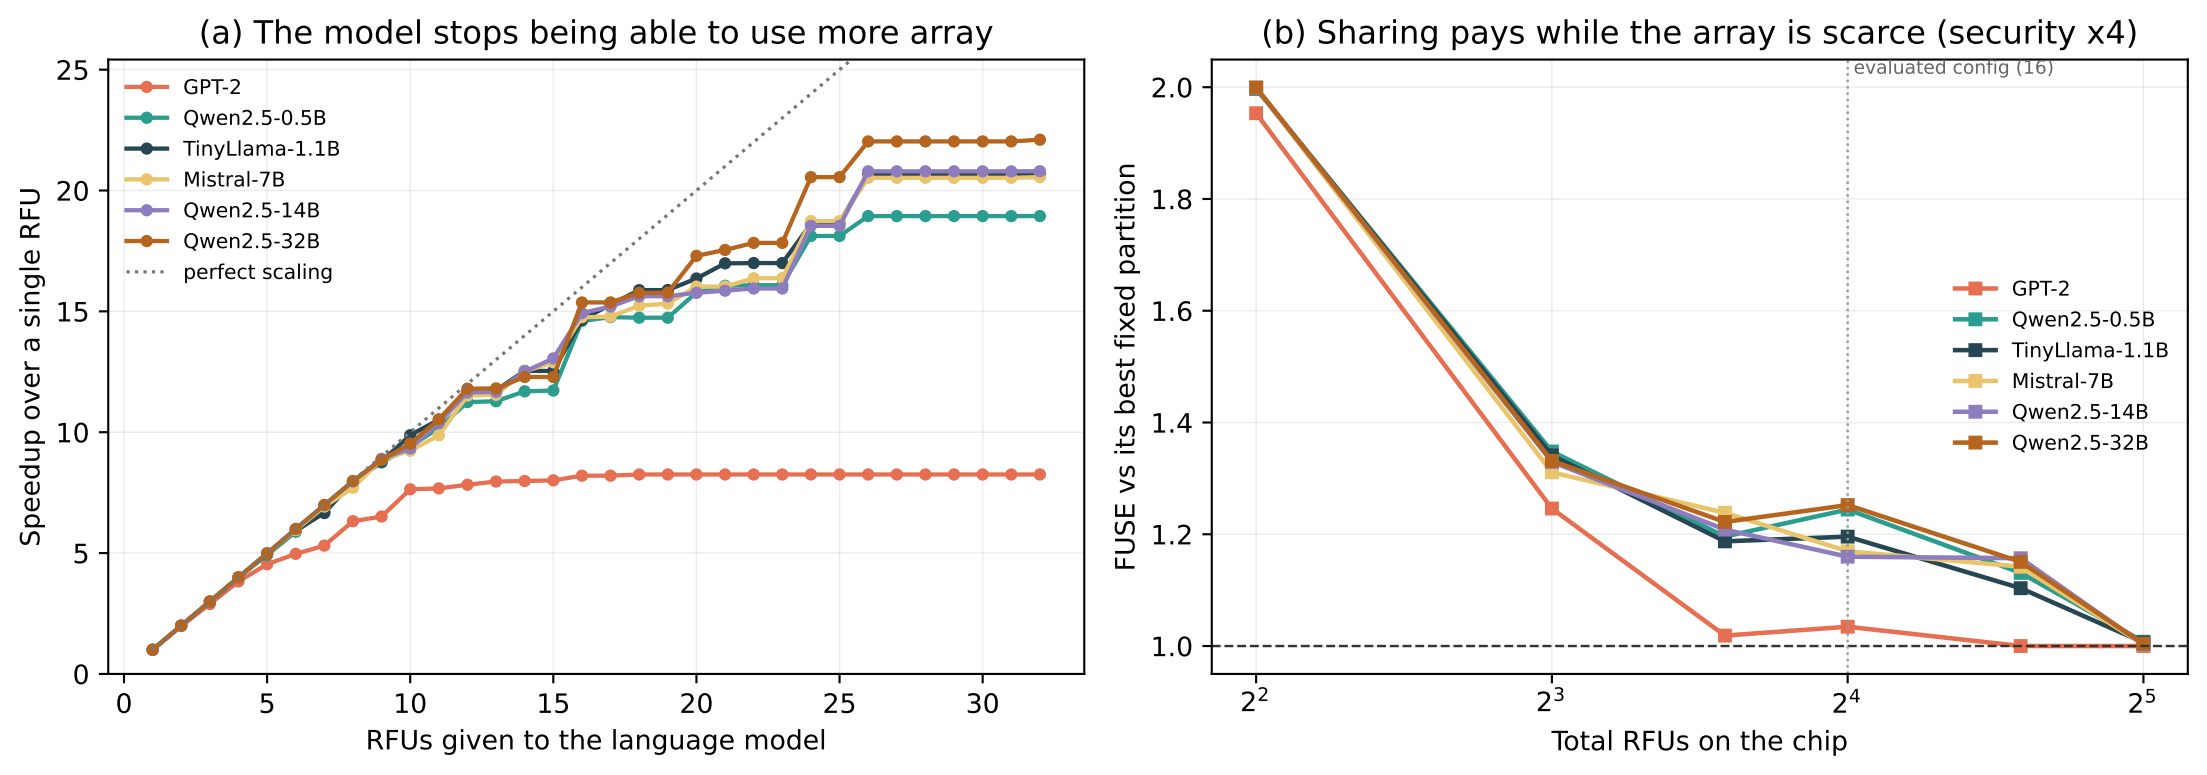
(a) The model stops being able to use more array
Speedup over a single RFU
RFUs given to the language model
GPT-2
Qwen2.5-0.5B
TinyLlama-1.1B
Mistral-7B
Qwen2.5-14B
Qwen2.5-32B
perfect scaling
0
5
10
15
20
25
30
(b) Sharing pays while the array is scarce (security x4)
FUSE vs its best fixed partition
Total RFUs on the chip
evaluated config (16)
1.0
1.2
1.4
1.6
1.8
2.0
2^2
2^3
2^4
2^5