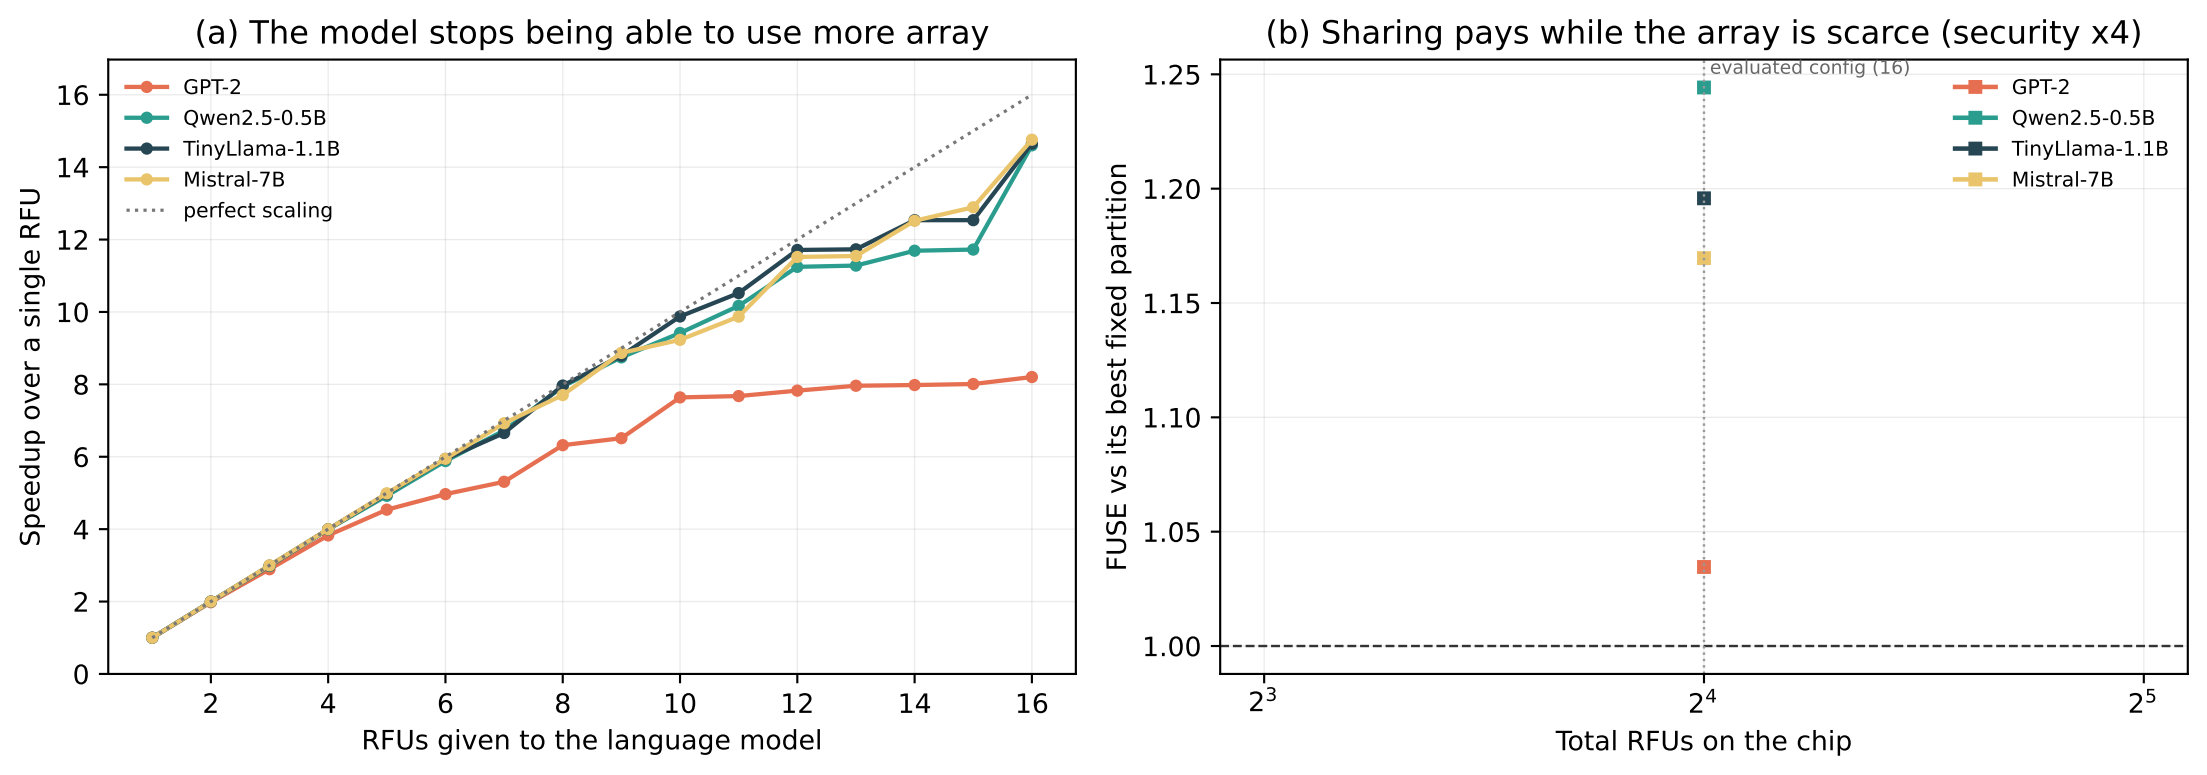
(a) The model stops being able to use more array
GPT-2
Qwen2.5-0.5B
TinyLlama-1.1B
Mistral-7B
perfect scaling
Speedup over a single RFU
RFUs given to the language model
0
2
4
6
8
10
12
14
16
(b) Sharing pays while the array is scarce (security x4)
evaluated config (16)
GPT-2
Qwen2.5-0.5B
TinyLlama-1.1B
Mistral-7B
FUSE vs its best fixed partition
1.00
1.05
1.10
1.15
1.20
1.25
2^3
2^4
2^5
Total RFUs on the chip