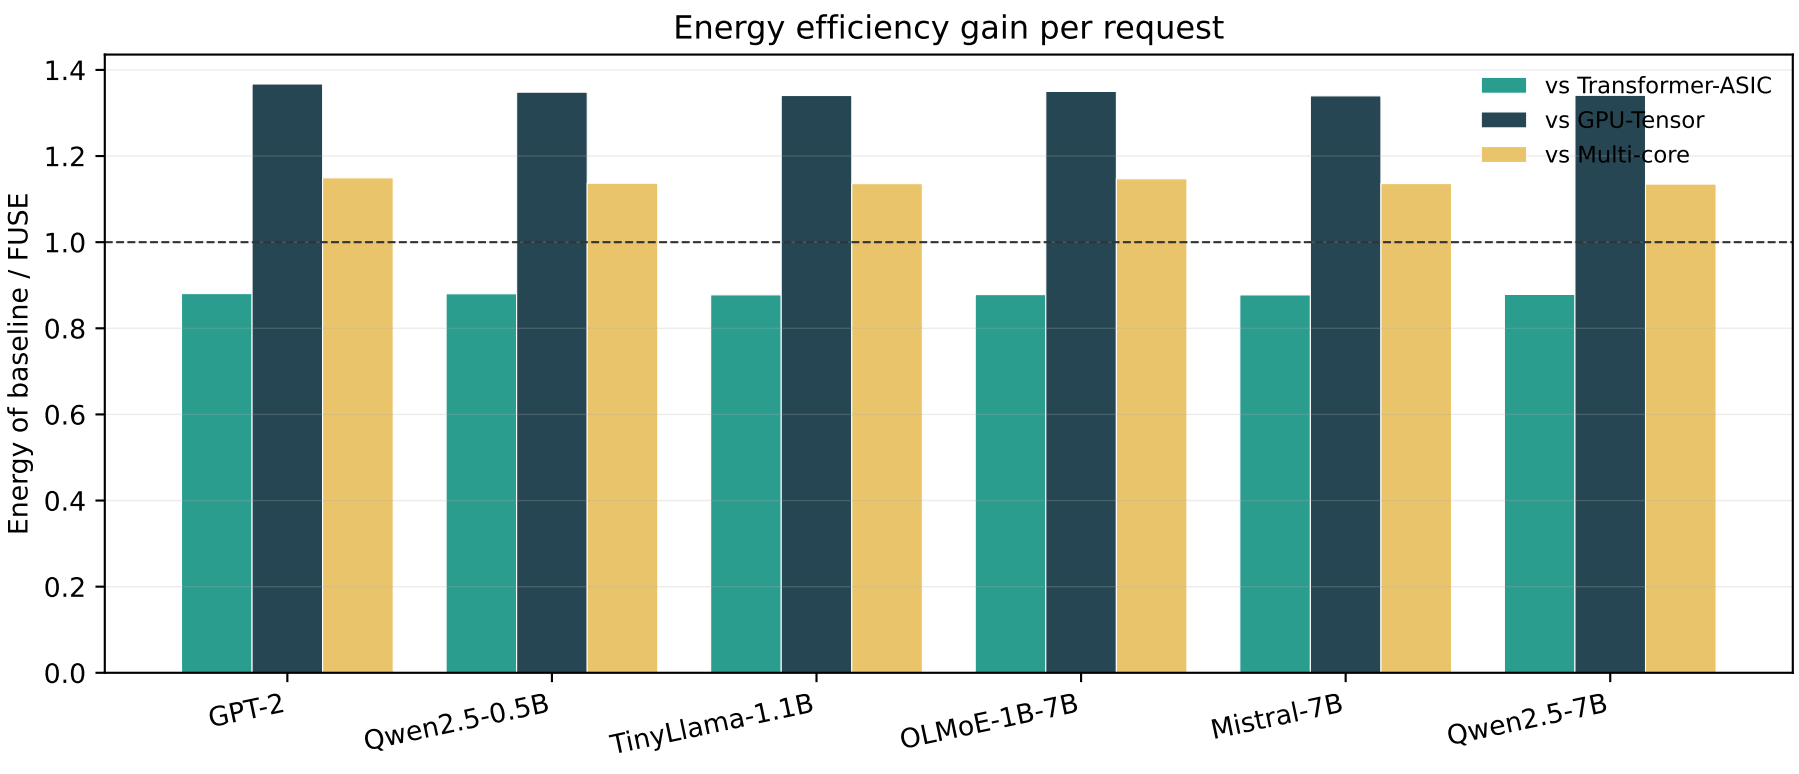
Energy efficiency gain per request
Energy of baseline / FUSE
1.4
1.2
1.0
0.8
0.6
0.4
0.2
0.0
vs Transformer-ASIC
vs GPU-Tensor
vs Multi-core
GPT-2
Qwen2.5-0.5B
TinyLlama-1.1B
OLMoE-1B-7B
Mistral-7B
Qwen2.5-7B