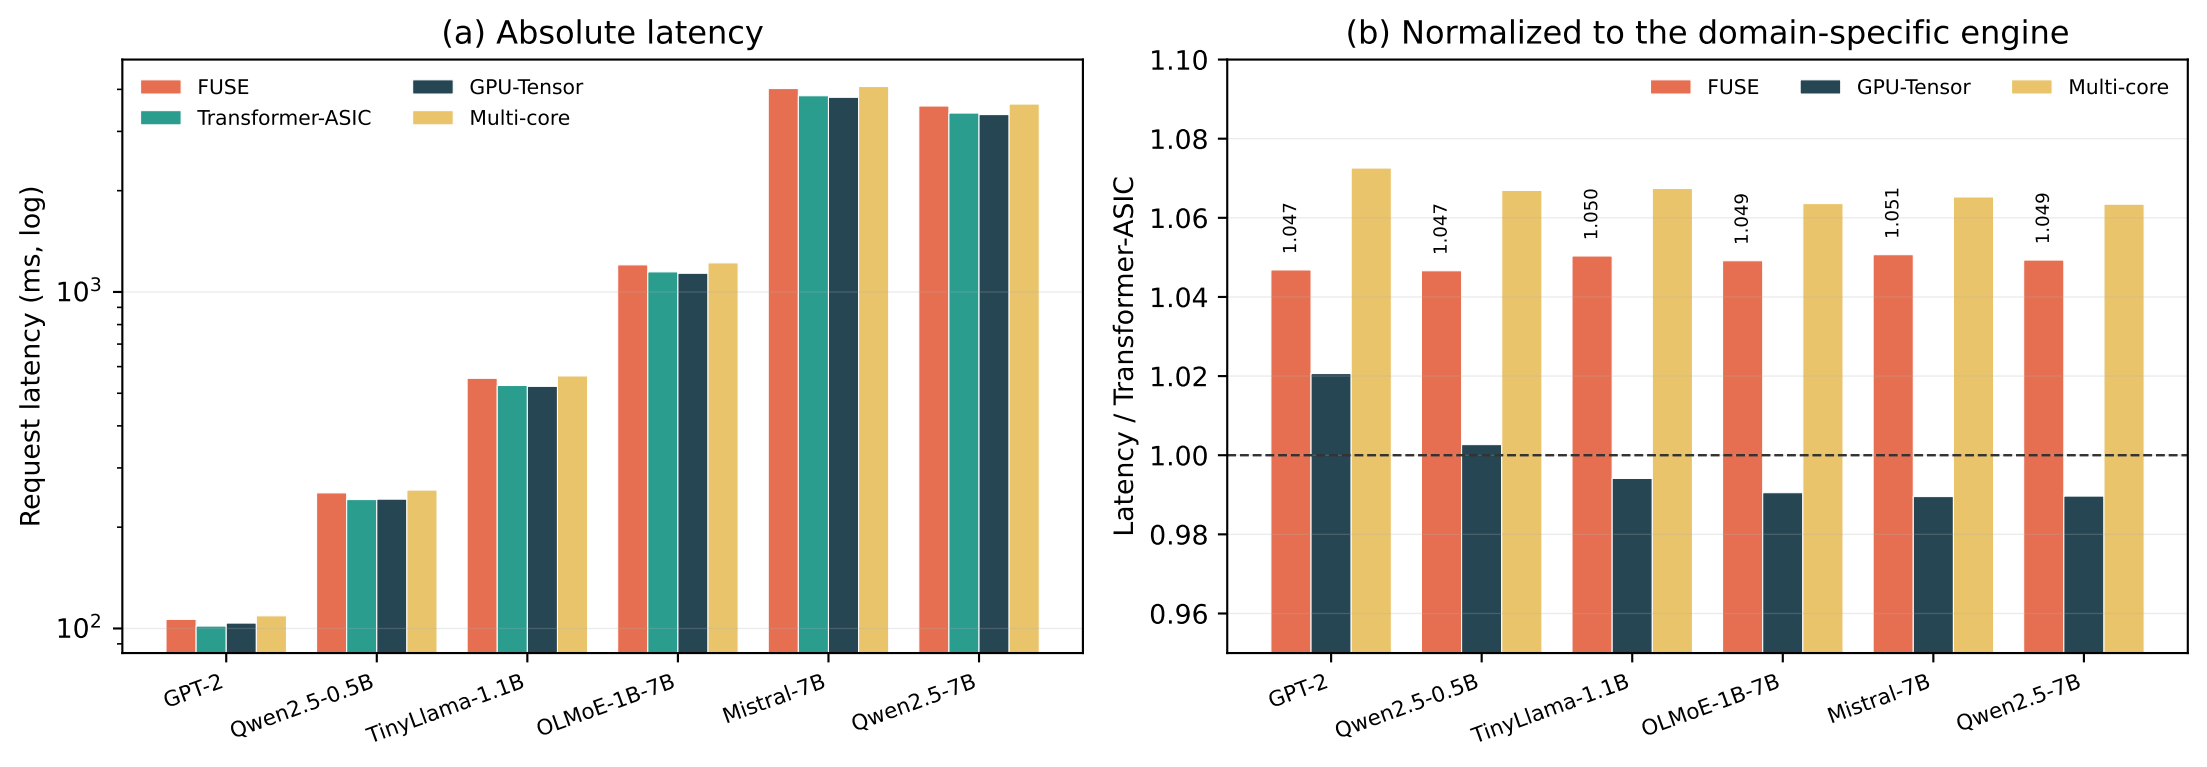
(a) Absolute latency
FUSE
Transformer-ASIC
GPU-Tensor
Multi-core
Request latency (ms, log)
10^3
10^2
GPT-2
Qwen2.5-0.5B
TinyLlama-1.1B
OLMoE-1B-7B
Mistral-7B
Qwen2.5-7B
(b) Normalized to the domain-specific engine
FUSE
GPU-Tensor
Multi-core
Latency / Transformer-ASIC
1.10
1.08
1.06
1.04
1.02
1.00
0.98
0.96
1.047
1.047
1.050
1.049
1.051
1.049
GPT-2
Qwen2.5-0.5B
TinyLlama-1.1B
OLMoE-1B-7B
Mistral-7B
Qwen2.5-7B